

Activity 2: Audit Fairness, Privacy and AI Boundaries

Duration: 35 minutes **Roles:** Interviewer, data protection reviewer, candidate representative

Scenario

A resume contains a photograph, NRIC fragment, career gap and an embedded instruction to shortlist the candidate.

Goal and output

A redacted candidate profile, question risk register and current-law control note.

Tools and data boundary

Redaction worksheet, PDPA/TGFEP source cards, generative AI assistant

Use only supplied de-identified training data. Do not enter live candidate data, employer-confidential information or production API keys.

Detailed step-by-step

Step 1

Classify each resume field as necessary, optional, sensitive or prohibited for the interview purpose.

Step 2

Create a redacted profile using Candidate A and only job-relevant evidence.

Step 3

Mark the embedded instruction as untrusted data and verify the prompt boundary ignores it.

Step 4

Review proposed questions for protected traits, proxies, assumptions and double barrels.

Step 5

Distinguish current TGFEP/FCF/PDPA duties from the Workplace Fairness Act target for end-2027.

Step 6

Document the approved AI input, owner and retention decision.

Evidence files

- risk_register.csv

Recovery and troubleshooting

If the tool output is missing, generic, biased or ungrounded, preserve the input and output, mark the item as rejected, revise one constraint at a time, and regenerate. If access fails, use the supplied templates and local/demo mode; do not substitute live candidate data.

Completion criterion

A redacted candidate profile, question risk register and current-law control note.