

Activity 4: Prompt Engineering: The Four-Part Pattern for Interview Questions

Duration: 40 minutes **Roles:** Question designer working with an independent reviewer

Scenario

A vague prompt produces generic questions. You will prove it, then fix it with a structured prompt.

Goal and output

A before/after prompt pair and a reviewed set of job-related questions.

Tools and data boundary

Microsoft 365 Copilot | the approved job description

Use only supplied de-identified training data. Do not enter live candidate data, employer-confidential information or production API keys.

Detailed step-by-step

Step 1

Run a deliberately weak prompt: 'Give me interview questions for a data analyst.' Save the output.

Step 2

Mark every defect in that output: generic, unanchored, leading, double-barrelled, or untraceable to a competency.

Step 3

Rewrite using the four-part pattern - ROLE, CONTEXT, TASK, CONSTRAINTS - naming the output format you want.

Step 4

Put the fairness rules in CONSTRAINTS: no protected traits, no proxies, one competency per question.

Step 5

Run the improved prompt and compare the two outputs side by side.

Step 6

Iterate once more: ask Copilot to critique its own questions against the evidence contract and revise.

Step 7

Record which specific words in your prompt caused the improvement.

Evidence files

- prompt_ab.txt
- question_review.csv

Recovery and troubleshooting

If the tool output is missing, generic, biased or ungrounded, preserve the input and output, mark the item as rejected, revise one constraint at a time, and regenerate. If access fails, use the supplied templates and local/demo mode; do not substitute live candidate data.

Completion criterion

A before/after prompt pair and a reviewed set of job-related questions.