

Activity 8: Score Evidence and Calibrate the Panel with Copilot

Duration: 40 minutes **Roles:** Three independent raters and a calibration chair

Scenario

Three raters disagree on the same candidate transcript. Copilot helps surface the evidence, but the panel decides.

Goal and output

Independent scores, a discrepancy discussion and a documented calibrated result.

Tools and data boundary

Evidence Scorer (KEEP) agent | your transcript from Activity 7

Use only supplied de-identified training data. Do not enter live candidate data, employer-confidential information or production API keys.

Detailed step-by-step

Step 1

Each rater scores the transcript independently against the anchors BEFORE any discussion or AI use.

Step 2

Paste the de-identified transcript into the Evidence Scorer agent and ask it to separate evidence from inference.

Step 3

Compare the agent's proposed scores with your own. Where you differ, identify whether the evidence or the anchor reading differs.

Step 4

Use N/E where evidence is genuinely missing. Confirm the agent did not invent evidence to fill a gap.

Step 5

Discuss any difference of two or more points by citing the transcript, never seniority.

Step 6

Record the final calibrated score and the rationale, and note that a human panel made the decision.

Step 7

Log any rubric wording that caused disagreement for the next revision.

Evidence files

- evidence_matrix.csv

Recovery and troubleshooting

If the tool output is missing, generic, biased or ungrounded, preserve the input and output, mark the item as rejected, revise one constraint at a time, and regenerate. If access fails, use the supplied templates and local/demo mode; do not substitute live candidate data.

Completion criterion

Independent scores, a discrepancy discussion and a documented calibrated result.